

# Supervised

2019-5-15

```
heart_disease = read_csv("./heart.csv") %>%
  mutate(target = ifelse(target==1, 0, 1)) %>%
  mutate(target=as.factor(target)) %>%
  mutate(target=as.factor(ifelse(target==0, "absence", "presence")))%>%
  mutate(target = relevel(target, "presence"))

heart_disease = heart_disease %>%
  filter(thal != 0) %>%
  mutate(sex=as.factor(sex),
         cp=as.factor(cp),
         fbs=as.factor(fbs),
         restecg=as.factor(restecg),
         exang=as.factor(exang),
         slope=as.factor(slope),
         thal=factor(thal))

model.x <- model.matrix(target~.,heart_disease)[,-1]
model.y <- heart_disease$target
```

## Regularized logistic

```
ctrl = trainControl(method = "cv",
                    classProbs = TRUE,
                    summaryFunction = twoClassSummary)


glmGrid <- expand.grid(.alpha = seq(0, 0.5, length = 10),
                     .lambda = exp(seq(-10,-1, length = 100)))

set.seed(1)
model.glm <- train(x = model.x,
                  y = model.y,
                  method = "glmnet",
                  tuneGrid = glmGrid,
                  metric = "ROC",
                  trControl = ctrl)

ggplot(model.glm, highlight = T) +
  viridis::scale_color_viridis(discrete = TRUE) +
  scale_shape_manual(values = seq(1,10))

## Scale for 'colour' is already present. Adding another scale for
## 'colour', which will replace the existing scale.

## Scale for 'shape' is already present. Adding another scale for 'shape',
## which will replace the existing scale.
```



```
model.glm$bestTune
```

```
## alpha lambda
## 95 0 0.2335065
```

```
glmnet = glmnet(x = model.x, y = model.y,
  family = "binomial",
  alpha = 0,
  lambda = 0.1946867)
broom::tidy(glmnet)
```

```
## # A tibble: 19 x 5
##   term      step estimate lambda dev.ratio
##   <chr>      <dbl>      <dbl> <dbl>      <dbl>
## 1 (Intercept) 1 0.624 0.195 0.431
## 2 age          1 -0.00885 0.195 0.431
## 3 sex1          1 -0.462 0.195 0.431
## 4 cp1           1 0.385 0.195 0.431
## 5 cp2           1 0.586 0.195 0.431
## 6 cp3           1 0.524 0.195 0.431
## 7 trestbps      1 -0.00476 0.195 0.431
## 8 chol          1 -0.00120 0.195 0.431
## 9 fbs1          1 0.0693 0.195 0.431
## 10 restecg1     1 0.246 0.195 0.431
## 11 restecg2     1 -0.198 0.195 0.431
## 12 thalach      1 0.00952 0.195 0.431
## 13 exang1       1 -0.522 0.195 0.431
## 14 oldpeak      1 -0.199 0.195 0.431
## 15 slope1       1 -0.293 0.195 0.431
## 16 slope2       1 0.290 0.195 0.431
## 17 ca           1 -0.306 0.195 0.431
## 18 thal2        1 0.527 0.195 0.431
## 19 thal3        1 -0.526 0.195 0.431
```

## LDA

```
set.seed(1)
model.lda = train(x = model.x,
                  y = model.y,
                  method = "lda",
                  metric = "ROC",
                  trControl = ctrl)
```

## Naive bayes


```
set.seed(1)
nbGrid = expand.grid(usekernel = c(FALSE,TRUE),
                    fL = 1, adjust = seq(0, 4, length = 20))
model.bayes = train(x = model.x,
                    y = model.y,
                    method = "nb",
                    tuneGrid = nbGrid,
                    metric = "ROC",
                    trControl = ctrl)

ggplot(model.bayes, highlight = T)

model.bayes$bestTune
```

##Tree

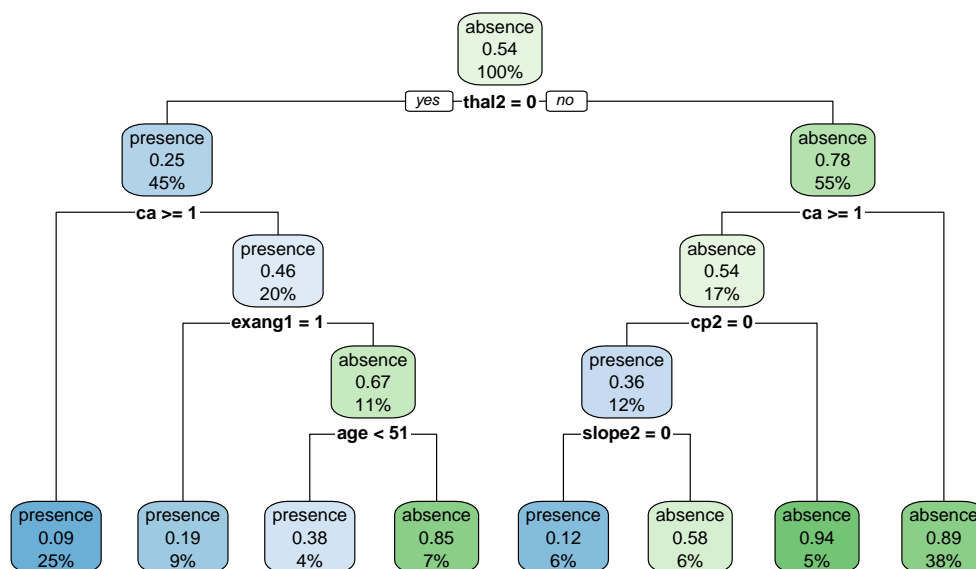
```
set.seed(1)
tree.class <- train(model.x, model.y,
                    method = "rpart",
                    tuneGrid = data.frame(cp = exp(seq(-10,-3, len = 20))),
                    trControl = ctrl,
                    metric = "ROC")
ggplot(tree.class, highlight = TRUE)
```



```
tree.class$bestTune
```

```
##           cp
## 13 0.003776539
```

```
rpart.plot(tree.class$finalModel)
```



```
##Bagging
```

```
bagging.grid <- expand.grid(mtry = 18,
  splitrule = "gini",
  min.node.size = 10:50)
```

```

set.seed(1)
bagging.class <- train(model.x, model.y,
  method = "ranger",
  tuneGrid = bagging.grid,
  metric = "ROC",
  trControl = ctrl,
  importance = "impurity")

ggplot(bagging.class, highlight = TRUE)
bagging.class$bestTune

barplot(sort(ranger::importance(bagging.class$finalModel),
  decreasing = FALSE),
  las = 2, horiz = TRUE, cex.names = 0.7,
  col = colorRampPalette(colors = c("darkred", "white", "darkblue"))(18))

```

##Random Forest

```

rf.grid <- expand.grid(mtry = 1:6,
  splitrule = "gini",
  min.node.size = seq(1,191, by = 2))

set.seed(1)
rf.class <- train(model.x, model.y,
  method = "ranger",
  tuneGrid = rf.grid,
  metric = "ROC",
  trControl = ctrl,
  importance = "impurity")

rf.class$bestTune

ggplot(rf.class, highlight = TRUE) +
  viridis::scale_color_viridis(discrete = TRUE) +
  scale_shape_manual(values = seq(1,7))

barplot(sort(ranger::importance(rf.class$finalModel), decreasing = FALSE),
  las = 2, horiz = TRUE, cex.names = 0.7,
  col = colorRampPalette(colors = c("darkred", "white", "darkblue"))(18))

```

##Boosting

```

boost.grid <- expand.grid(n.trees = seq(20, 1700, by = 25),
  interaction.depth = 1:6,
  shrinkage = seq(0.005, 0.06, by = 0.005),
  n.minobsinnode = 1)

set.seed(1)
# Adaboost loss function
boost.class = train(model.x, model.y,
  tuneGrid = boost.grid,
  trControl = ctrl,
  method = "gbm",
  distribution = "adaboost",

```

```

        metric = "ROC",
        verbose = FALSE)

boost.class$bestTune

ggplot(boost.class, highlight = TRUE) +
  viridis::scale_color_viridis(discrete = TRUE) +
  scale_shape_manual(values = seq(0,11))

summary(boost.class$finalModel, las = 2, cBars = 19, cex.names = 0.6)

```


## SVM ROC

```

## linear boundary
set.seed(1)
svml.fit <- train(target~.,
  data = heart_disease,
  method = "svmLinear2",
  preProcess = c("center", "scale"),
  tuneGrid = data.frame(cost = exp(seq(-7,-2,len=50))),
  trControl = ctrl,
  metric = "ROC")

ggplot(svml.fit, highlight = TRUE)

```



```

svml.fit$bestTune

##          cost
## 11 0.002529859

## radial kernel
svmr.grid <- expand.grid(C = exp(seq(-4,5,len=50)),

```

```

                                sigma = exp(seq(-5,-2,len=10)))
set.seed(1)
svmr.fit <- train(target~.,
                  data = heart_disease,
                  method = "svmRadial",
                  preProcess = c("center", "scale"),
                  tuneGrid = svmr.grid,
                  trControl = ctrl,
                  metric = "ROC")

ggplot(svmr.fit, highlight = TRUE) +
  viridis::scale_color_viridis(discrete = TRUE) +
  scale_shape_manual(values = seq(1,10))
svmr.fit$bestTune

```

## Neural network

```

nnetGrid <- expand.grid(size = seq(from = 16, to = 30, by = 2),
                       decay = seq(from = 5, to = 8, length = 30))

set.seed(1)
cnnet.fit <- train(target~.,
                  heart_disease,
                  method = "nnet",
                  tuneGrid = nnetGrid,
                  preProcess = c("center", "scale"),
                  trControl = ctrl,
                  metric = "ROC",
                  trace = FALSE)

ggplot(cnnet.fit, highlight = TRUE) +
  viridis::scale_color_viridis(discrete = TRUE) +
  scale_shape_manual(values = seq(1,13))

cnnet.fit$bestTune

```

```

load(file = "./saved_results/cnnet.rda")
load(file = "./saved_results/boost.rda")
load(file = "./saved_results/rf.rda")
load(file = "./saved_results/bagging.rda")
load(file = "./saved_results/bayes.rda")
load(file = "./saved_results/svmr.rda")

resamp = resamples(list(
  Regularized_logistic = model.glm,
  LDA = model.lda,
  Naive_Bayes = model.bayes,
  Adaboost = boost.class,
  Random_forests = rf.class,
  Bagging = bagging.class,
  Tree = tree.class,
  Neural_network = cnnet.fit,
  SVM_linear = svmr.fit,

```

```

SVM_gaussian = svmr.fit
))
summary(resamp)


##
## Call:
## summary.resamples(object = resamp)
##
## Models: Regularized_logistic, LDA, Naive_Bayes, Adaboost, Random_forests, Bagging, Tree, Neural_netw
## Number of resamples: 10
##
## ROC
##
##           Min.   1st Qu.   Median     Mean   3rd Qu.
## Regularized_logistic 0.8125000 0.8612839 0.8998162 0.9021715 0.9553571
## LDA                  0.7812500 0.8463660 0.8998162 0.8981719 0.9497768
## Naive_Bayes          0.8169643 0.8725103 0.8994829 0.9097952 0.9681490
## Adaboost             0.7678571 0.8918572 0.9067752 0.9062419 0.9434086
## Random_forests       0.8125000 0.8747424 0.9106335 0.9091185 0.9637605
## Bagging              0.7901786 0.8613445 0.8943924 0.8934995 0.9357224
## Tree                 0.7626050 0.7968750 0.8471386 0.8446792 0.8771008
## Neural_network       0.8125000 0.8641827 0.9010989 0.9026261 0.9553571
## SVM_linear           0.8109244 0.8605769 0.9056238 0.9043815 0.9658310
## SVM_gaussian         0.8214286 0.8567590 0.9012605 0.9027614 0.9497768
##
##           Max. NA's
## Regularized_logistic 0.9747899    0
## LDA                  0.9903846    0
## Naive_Bayes          0.9821429    0
## Adaboost             0.9807692    0
## Random_forests       0.9776786    0
## Bagging              0.9776786    0
## Tree                 0.9665179    0
## Neural_network       0.9747899    0
## SVM_linear           0.9789916    0
## SVM_gaussian         0.9873950    0
##
## Sens
##
##           Min.   1st Qu.   Median     Mean   3rd Qu.
## Regularized_logistic 0.5714286 0.6923077 0.7857143 0.7950549 0.9065934
## LDA                  0.5714286 0.6978022 0.7857143 0.7725275 0.8310440
## Naive_Bayes          0.6428571 0.7857143 0.8159341 0.8258242 0.9038462
## Adaboost             0.6153846 0.6978022 0.7857143 0.7879121 0.8571429
## Random_forests       0.6428571 0.6923077 0.7142857 0.7659341 0.8310440
## Bagging              0.6428571 0.6978022 0.7500000 0.7659341 0.8310440
## Tree                 0.6153846 0.7280220 0.7774725 0.7725275 0.8392857
## Neural_network       0.5714286 0.6923077 0.7857143 0.7950549 0.9065934
## SVM_linear           0.5714286 0.6552198 0.7500000 0.7653846 0.8543956
## SVM_gaussian         0.5384615 0.6401099 0.7857143 0.7653846 0.8571429
##
##           Max. NA's
## Regularized_logistic 1.0000000    0
## LDA                  1.0000000    0
## Naive_Bayes          1.0000000    0
## Adaboost             1.0000000    0
## Random_forests       1.0000000    0
## Bagging              0.9285714    0

```



```
## Tree          0.9285714    0
## Neural_network 1.0000000    0
## SVM_linear     1.0000000    0
## SVM_gaussian  1.0000000    0
##
## Spec
##              Min.   1st Qu.   Median     Mean   3rd Qu.
## Regularized_logistic 0.7500000 0.8152574 0.8786765 0.8838235 0.9411765
## LDA                  0.7500000 0.8152574 0.9099265 0.8838235 0.9411765
## Naive_Bayes          0.7500000 0.7500000 0.8235294 0.8470588 0.9264706
## Adaboost             0.6875000 0.8281250 0.9375000 0.8775735 0.9402574
## Random_forests       0.7058824 0.7812500 0.8786765 0.8658088 0.9402574
## Bagging              0.6875000 0.7766544 0.8235294 0.8297794 0.8621324
## Tree                 0.5625000 0.8152574 0.8492647 0.8349265 0.8823529
## Neural_network       0.7500000 0.8152574 0.8786765 0.8838235 0.9411765
## SVM_linear           0.7500000 0.7766544 0.8786765 0.8658088 0.9264706
## SVM_gaussian         0.7500000 0.8125000 0.8492647 0.8536765 0.9237132
##
##              Max. NA's
## Regularized_logistic 1.0000000    0
## LDA                  1.0000000    0
## Naive_Bayes          1.0000000    0
## Adaboost             0.9411765    0
## Random_forests       1.0000000    0
## Bagging              1.0000000    0
## Tree                 0.9375000    0
## Neural_network       1.0000000    0
## SVM_linear           1.0000000    0
## SVM_gaussian         0.9411765    0
```

```
bwplot(resamp, metric = "ROC")
```



```
###centered ICE
```

```

ice_thalach.rf = rf.class %>%
  pdp::partial(pred.var = "thalach",
    grid.resolution = 100,
    ice = TRUE,
    prob = TRUE) %>%
  autoplot(train = heart_disease, alpha = .1) +
  ggtitle("Random forest, thalach")

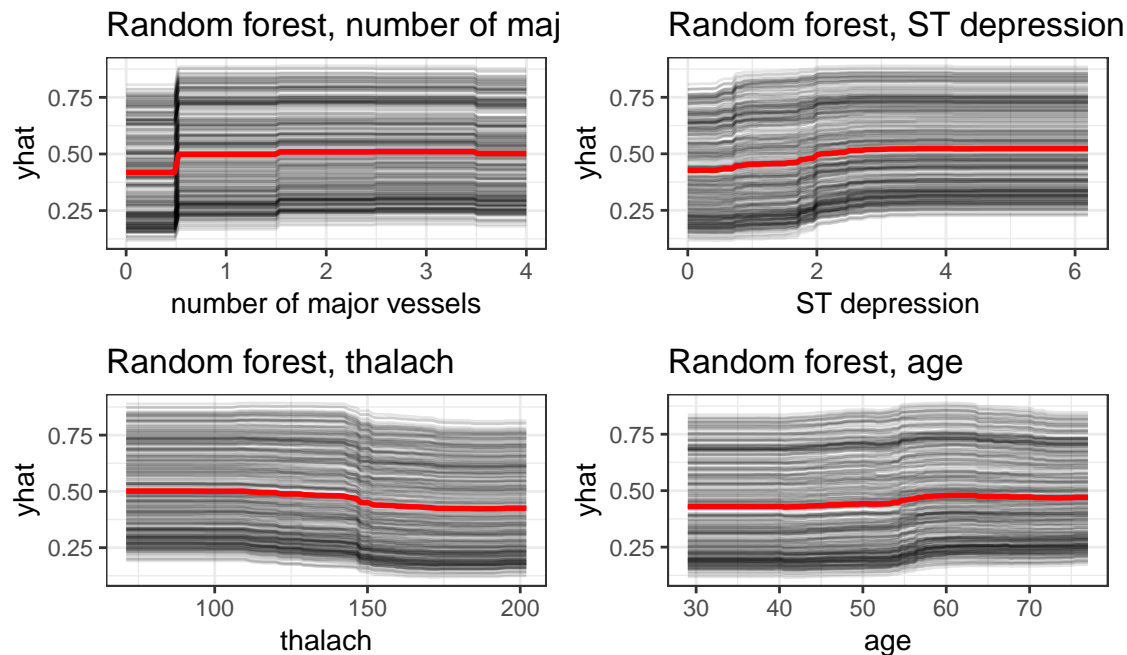
ice_ca.rf = rf.class %>%
  pdp::partial(pred.var = "ca",
    grid.resolution = 100,
    ice = TRUE,
    prob = TRUE) %>%
  autoplot(train = heart_disease, alpha = .1,
    xlab = "number of major vessels") +
  ggtitle("Random forest, number of major vessels")

ice_oldpeak.rf = rf.class %>%
  partial(pred.var = "oldpeak",
    grid.resolution = 100,
    ice = TRUE,
    prob = TRUE) %>%
  autoplot(train = heart_disease, alpha = .1,
    xlab = "ST depression") +
  ggtitle("Random forest, ST depression")

ice_age.rf = rf.class %>%
  pdp::partial(pred.var = "age",
    grid.resolution = 100,
    ice = TRUE,
    prob = TRUE) %>%
  autoplot(train = heart_disease, alpha = .1) +
  ggtitle("Random forest, age")

grid.arrange(ice_ca.rf, ice_oldpeak.rf,
  ice_thalach.rf, ice_age.rf, nrow = 2)

```



## Variable importance

```
library(gbm)
```

```
## Loaded gbm 2.1.5
```

```
varImp(model.glm)
```

```
## glmnet variable importance
```

```
##
```

```
## Overall
```

```
## cp2 100.0000
```

```
## thal2 94.8188
```

```
## thal3 93.4880
```

```
## exang1 92.3686
```

```
## cp3 87.3424
```

```
## sex1 79.4864
```

```
## cp1 66.9913
```

```
## slope2 52.7207
```

```
## slope1 52.7011
```

```
## ca 52.4277
```

```
## restecg1 42.4609
```

```
## restecg2 36.5417
```

```
## oldpeak 34.8283
```

```
## fbs1 10.6763
```

```
## thalach 1.4898
```

```
## age 1.4370
```

```
## trestbps 0.6164
```

```
## chol 0.0000
```

```
varImp(model.lda)
```

```
## ROC curve variable importance
```

```
##
##          Importance
## thal2      100.0000
## thalach     95.2812
## thal3      90.1317
## oldpeak     89.9165
## ca          89.7356
## exang1      76.2594
## slope2      73.0991
## slope1      66.7700
## cp2         52.2776
## age         51.1237
## sex1        47.3435
## cp1         33.1353
## restecg1    33.0578
## chol        25.1356
## trestbps    24.9634
## cp3         6.4669
## restecg2     0.5339
## fbs1        0.0000
```

```
varImp(model.bayes)
```

```
## ROC curve variable importance
##
##          Importance
## thal2      100.0000
## thalach     95.2812
## thal3      90.1317
## oldpeak     89.9165
## ca          89.7356
## exang1      76.2594
## slope2      73.0991
## slope1      66.7700
## cp2         52.2776
## age         51.1237
## sex1        47.3435
## cp1         33.1353
## restecg1    33.0578
## chol        25.1356
## trestbps    24.9634
## cp3         6.4669
## restecg2     0.5339
## fbs1        0.0000
```

```
varImp(boost.class)
```

```
## gbm variable importance
##
##          Overall
## ca          100.0000
## oldpeak     79.6457
## thal2       77.0270
## thalach     61.2231
## chol        52.8994
```

```
## trestbps 50.8206
## age 42.3494
## exang1 40.8228
## cp2 28.2991
## thal3 28.0742
## sex1 19.9508
## cp3 19.4944
## slope2 17.4076
## slope1 9.4854
## restecg1 8.0607
## cp1 7.8021
## fbs1 0.9966
## restecg2 0.0000
```

```
varImp(rf.class)
```

```
## ranger variable importance
##
## Overall
## thal2 100.000
## ca 94.184
## thal3 78.793
## oldpeak 77.081
## thalach 71.046
## exang1 59.586
## slope2 49.347
## age 48.026
## cp2 34.566
## chol 33.949
## sex1 32.134
## slope1 31.905
## trestbps 30.498
## cp1 15.616
## restecg1 10.725
## cp3 8.663
## fbs1 2.118
## restecg2 0.000
```

```
varImp(bagging.class)
```

```
## ranger variable importance
##
## Overall
## thal2 100.00000
## ca 47.18886
## oldpeak 32.70738
## thalach 27.53944
## age 14.77607
## exang1 13.43368
## trestbps 10.65010
## cp2 10.41997
## thal3 9.60369
## chol 8.71625
## cp3 5.61067
## slope2 3.96182
```

```
## sex1      3.24090
## restecg1  1.96555
## slope1    1.74960
## cp1       0.74229
## fbs1      0.04796
## restecg2  0.00000

varImp(tree.class)

## rpart variable importance
##
##           Overall
## thalach  100.000
## ca       96.684
## exang1    87.087
## thal2     80.067
## thal3     68.075
## oldpeak   37.514
## age       25.951
## slope2    17.578
## cp2       15.249
## sex1      14.902
## slope1    13.928
## cp3        8.264
## chol       7.543
## cp1        4.269
## fbs1       0.000
## restecg1   0.000
## restecg2   0.000
## trestbps   0.000
```

```
varImp(cnnnet.fit)

## nnet variable importance
##
##           Overall
## ca       100.000
## thal2     85.381
## thal3     82.923
## cp2       82.012
## exang1    77.280
## oldpeak   71.835
## thalach   68.244
## sex1      65.485
## slope2    44.568
## slope1    44.370
## cp1       40.316
## cp3       38.814
## restecg1  32.957
## trestbps  21.097
## age       21.068
## chol      12.338
## fbs1       1.239
## restecg2   0.000
```

```
#Comparing accuracy
```

##Regularized logistic

```
ctrl2 <- trainControl(method = "cv")

glmGrid <- expand.grid(.alpha = 0,
                      .lambda = 0.2335065)

set.seed(1)
model.glm.2 <- train(x = model.x,
                    y = model.y,
                    tuneGrid = glmGrid,
                    method = "glmnet",
                    metric = "Accuracy",
                    trControl = ctrl2)
```

##LDA

```
set.seed(1)
model.lda.2 = train(x = model.x,
                   y = model.y,
                   method = "lda",
                   metric = "Accuracy",
                   trControl = ctrl2)
```

##Naive bayes

```
set.seed(1)
nbGrid = expand.grid(usekernel = TRUE,
                    fL = 1, adjust = 1.473684)
model.bayes.2 = train(x = model.x,
                    y = model.y,
                    method = "nb",
                    tuneGrid = nbGrid,
                    metric = "Accuracy",
                    trControl = ctrl2)
```

##Tree

```
set.seed(1)
tree.class.2 <- train(model.x, model.y,
                    method = "rpart",
                    tuneGrid = data.frame(cp = 0.003776539),
                    trControl = ctrl2,
                    metric = "Accuracy")
```

##Bagging

```
bagging.grid <- expand.grid(mtry = 18,
                          splitrule = "gini",
                          min.node.size = 40)

set.seed(1)
bagging.class.2 <- train(model.x, model.y,
                      method = "ranger",
                      tuneGrid = bagging.grid,
                      metric = "Accuracy",
                      trControl = ctrl2,
                      importance = "impurity")
```

```
##Random Forest
```

```
rf.grid <- expand.grid(mtry = 1,  
                      splitrule = "gini",  
                      min.node.size = 25)
```

```
set.seed(1)
```

```
rf.class.2 <- train(model.x, model.y,  
                   method = "ranger",  
                   tuneGrid = rf.grid,  
                   metric = "Accuracy",  
                   trControl = ctrl2,  
                   importance = "impurity")
```

```
##Boosting
```

```
boost.grid <- expand.grid(n.trees = 1370,  
                        interaction.depth = 1,  
                        shrinkage = 0.015,  
                        n.minobsinnode = 1)
```

```
set.seed(1)
```

```
# Adaboost loss function
```

```
boost.class.2 = train(model.x, model.y,  
                    tuneGrid = boost.grid,  
                    trControl = ctrl2,  
                    method = "gbm",  
                    distribution = "adaboost",  
                    metric = "Accuracy",  
                    verbose = FALSE)
```

## Neural network

```
nnetGrid <- expand.grid(size = 18,  
                      decay = 6.448276)
```

```
set.seed(1)
```

```
cnnet.fit.2 <- train(target~.,  
                   heart_disease,  
                   method = "nnet",  
                   tuneGrid = nnetGrid,  
                   preProcess = c("center", "scale"),  
                   trControl = ctrl2,  
                   metric = "Accuracy",  
                   trace = FALSE)
```

## SVM

```
## linear boundary
```


```
set.seed(1)
```

```
svml.fit.2 <- train(target~.,  
                  data = heart_disease,  
                  method = "svmLinear2",
```



```
preProcess = c("center", "scale"),
tuneGrid = data.frame(cost = exp(seq(-7,-2,len=50))),
trControl = ctrl12)
```

```
ggplot(svm1.fit, highlight = TRUE)
```




```
svm1.fit$bestTune
```

```
##          cost
## 11 0.002529859
```

```
## radial kernel
svmr.grid <- expand.grid(C = exp(seq(-4,5,len=50)),
                        sigma = exp(seq(-5,-2,len=10)))
```

```
set.seed(1)
svmr.fit.2 <- train(target~.,
                    data = heart_disease,
                    method = "svmRadial",
                    preProcess = c("center", "scale"),
                    tuneGrid = svmr.grid,
                    trControl = ctrl12)
```

```
ggplot(svmr.fit, highlight = TRUE) +
  viridis::scale_color_viridis(discrete = TRUE) +
  scale_shape_manual(values = seq(1,10))
```



sigma

- 0.006737947
- 0.013123729
- 0.025561533
- 0.049787068
- 0.09697
- 0.009403563
- 0.018315639
- 0.035673993
- 0.069483451
- 0.13533

```
svmr.fit$bestTune
```

```
##          sigma      C
## 373 0.01312373 16.37766
```

```
resamp = resamples(list(
  glm.fit = model.glm.2,
  lda.fit = model.lda.2,
  bayes.fit = model.bayes.2,
  boost = boost.class.2,
  rf = rf.class.2,
  bagging = bagging.class.2,
  tree = tree.class.2,
  cnnet.fit = cnnet.fit.2,
  svml.fit = svml.fit.2,
  svmr.fit = svmr.fit.2
))
```

```
summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: glm.fit, lda.fit, bayes.fit, boost, rf, bagging, tree, cnnet.fit, svml.fit, svmr.fit
## Number of resamples: 10
##
## Accuracy
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## glm.fit  0.7000000 0.8031609 0.8360215 0.8441416 0.9000000 0.9677419    0
## lda.fit  0.7096774 0.7732759 0.8333333 0.8340267 0.8846774 0.9677419    0
## bayes.fit 0.7000000 0.7732759 0.8526882 0.8374750 0.8916667 0.9677419    0
## boost    0.7096774 0.7482759 0.8500000 0.8341416 0.8927419 0.9655172    0
## rf       0.6774194 0.7606322 0.8032258 0.8175677 0.9066092 0.9354839    0
```

```
## bagging 0.6666667 0.7806452 0.8331479 0.8141268 0.8562291 0.9000000 0
## tree 0.6666667 0.7789210 0.8166667 0.8066704 0.8666667 0.8709677 0
## cnet.fit 0.7000000 0.8031609 0.8360215 0.8441416 0.9000000 0.9677419 0
## svmf.fit 0.6774194 0.8000000 0.8360215 0.8378124 0.9155172 0.9677419 0
## svmr.fit 0.7419355 0.8068966 0.8500000 0.8473674 0.8927419 0.9354839 0
##
## Kappa
## Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## glm.fit 0.3946188 0.5944980 0.6694856 0.6826138 0.8004166 0.9352818 0
## lda.fit 0.4025696 0.5423267 0.6603832 0.6617627 0.7653612 0.9352818 0
## bayes.fit 0.3946188 0.5469194 0.7004056 0.6724875 0.7814956 0.9352818 0
## boost 0.4101480 0.4940966 0.6916528 0.6626159 0.7826973 0.9307876 0
## rf 0.3404255 0.5082084 0.6012348 0.6283910 0.8082234 0.8697479 0
## bagging 0.3303571 0.5478123 0.6580195 0.6220325 0.7099677 0.7963801 0
## tree 0.3421053 0.5469316 0.6266968 0.6090815 0.7315396 0.7427386 0
## cnet.fit 0.3946188 0.5944980 0.6694856 0.6826138 0.8004166 0.9352818 0
## svmf.fit 0.3404255 0.5893826 0.6709540 0.6710978 0.8279169 0.9352818 0
## svmr.fit 0.4655172 0.6095944 0.6904463 0.6889089 0.7861955 0.8697479 0
```

```
bwplot(resamp)
```

